

1

a.

$$N_{\text{page}}(\text{Parts}) = N_{\text{Tuple}}(\text{Parts}) / \text{TupleperPage}$$

$$= 60000 / 50$$

$$= 1200$$

$$N_{\text{page}}(\text{Supply}) = N_{\text{Tuple}}(\text{Supply}) / \text{TupleperPage}$$

$$= 150000 / 50$$

$$= 3000$$

$$\text{Cost} = N_{\text{Pages}}(\text{Parts}) + N_{\text{Pages}}(\text{Parts}) \times N_{\text{Pages}}(\text{Supply})$$

$$= 1200 + 1200 \times 3000 = \mathbf{3601200 \text{ I/O}}$$

b.

$$N_{\text{Blocks}}(\text{Parts}) = \text{ceil}(N_{\text{Pages}}(\text{Parts}) / (\text{Buffer} - 2)) = \text{ceil}(1200 / 200) = 6$$

$$\text{Cost} = N_{\text{Pages}}(\text{Parts}) + N_{\text{Blocks}}(\text{Parts}) \times N_{\text{Pages}}(\text{Supply})$$

$$= 1200 + 6 \times 3000 = \mathbf{19200 \text{ I/O}}$$

c.

$$\text{Cost} = \text{Sort}(\text{Parts}) + \text{Sort}(\text{Supply}) + \text{Merge}$$

$$= 2 \times \text{NumPasses} \times N_{\text{Pages}}(\text{Parts}) + (2 \times \text{NumPasses} \times N_{\text{Pages}}(\text{Supply}))$$

$$+ N_{\text{Pages}}(\text{Parts}) + N_{\text{Pages}}(\text{Supply})$$

$$= 5 \times (1200 + 3000) = \mathbf{21000 \text{ I/O}}$$

d.

$$\text{Cost} = 3 \times N_{\text{Pages}}(\text{Parts}) + 3 \times N_{\text{Pages}}(\text{Supply})$$

$$= 3 \times (1200 + 3000) = \mathbf{12600 \text{ I/O}}$$

e.

Block Nested Loops Join with B chosen so that the smaller table fits into memory as a single block.

$$N\text{Pages}(\text{Parts})/(\text{Buffer} - 2) = 1 \Rightarrow \text{Buffer} = \mathbf{1202 \text{ pages}}$$

$$\text{Cost} = 1200 + 3000 = \mathbf{4200 \text{ I/O}}$$

2

a.

$$\text{RF}(\text{Salary} > 300000) = (\text{High}(\text{Salary}) - \text{Value}) / (\text{High}(\text{Salary}) - \text{Low}(\text{Salary})) \\ = (500000 - 300000) / (500000 - 100000) = 1/2$$

$$\text{RF}(\text{Department} = \text{'Marketing'}) = 1 / \text{NKeys}(\text{Department}) = 1/6$$

$$\text{Result size} = \text{NTuples}(\text{Employee}) \times \pi_{\text{RF}}$$

$$= \text{NPages}(\text{Employee}) \times \text{NTuplesPerPage}(\text{Employee}) \times \text{RF}(\text{Salary} > 300000) \times \text{RF}(\text{Department} = \text{'Marketing'})$$

$$= 1200 \times 120 \times 1/2 \times 1/6 = \mathbf{12000 \text{ tuples}}$$

b.

Using clustered B+ tree on (Department, Salary):

$$\text{Cost} = \text{RF}(\text{Department} = \text{'Marketing'}) \times \text{RF}(\text{Salary} > 300000) \times \\ (\text{NPages}(\text{Index}) + \text{NPages}(\text{Employee}))$$

$$= 1/6 \times 1/2 \times (300 + 1200) = 125 \text{ I/O}$$

Using full table scan:

$$\text{Cost} = \text{NPages}(\text{Employee}) = 1200 \text{ I/O}$$

The best plan is the clustered B+ tree on (Department, Salary) with a cost of **125 I/O**.

c.

Using unclustered B+ tree on (Salary):

$$\text{Cost} = \text{RF}(\text{salary} > 300000) \times (\text{NPages}(\text{Index}) + \text{NTuples}(\text{Employee})) \\ = 1/2 \times (200 + 1200 \times 120) = 72100 \text{ I/O}$$

$$\text{Using full table scan: Cost} = \text{NPages}(\text{Employee}) = 1200 \text{ I/O}$$

The best plan is the full table scan with a cost of **1200 I/O**.

d.

Using unclustered hash on (Department):

$$\begin{aligned}\text{Cost} &= \text{RF}(\text{department} = \text{'Marketing'}) \times 2.2 \times \text{NTuples}(\text{Employee}) \\ &= 1/6 \times 2.2 \times 1200 \times 120 = 52800 \text{ I/O}\end{aligned}$$

Using full table scan:

$$\text{Cost} = \text{NPages}(\text{Employee}) = 1200 \text{ I/O}$$

The best plan is the full table scan with a cost of **1200 I/O**.

e.

Hash index cannot be used for range queries.

The only available plan is the full table scan with a cost of 1200 I/O.

3

a.

$$R_{Feid} = 1 / NKeys(eid) = 1 / NTuples(Employee) = 1 / 5000$$

$$R_{Fprojid} = 1 / NKeys(projid) = 1 / NTuples(Project) = 1 / 60000$$

$$R_{F(salary < 300000)} = (Value - Low(salary)) / (High(salary) - Low(salary)) = (300000 - 100000) / (500000 - 100000) = 1/2$$

$$R_{F(code = 'alpha 340')} = 1 / NKeys(code) = 1 / 1000$$

$$Result\ size = NTuples(Employee) \times NTuples(Project) \times$$

$$NTuples(Department) \times R_{Feid} \times R_{Fprojid} \times R_{F(salary < 300000)} \times R_{F(code = 'alpha 340')}$$

$$= 5000 \times 60000 \times 20000 \times 1/5000 \times 1/60000 \times 1/2 \times 1/1000$$

$$= \mathbf{10\ tuples}$$

b.(plan1)

$$NPages(Employee) = 5000 / 100 = 50\ pages$$

$$NPages(Project) = 60000 / 100 = 600\ pages$$

$$NPages(Department) = 20000 / 100 = 200\ pages$$

$$Cost\ to\ join\ Employee \bowtie Project = NPages(Employee) + NPages(Employee)$$

$$\times NPages(Project)$$

$$= 50 + 50 \times 600 = 30050\ I/O$$

$$Result\ size\ of\ Employee \bowtie Project = NTuples(Employee) \times NTuples(Project)$$

$$\times 1/NKeys(eid)$$

$$= 5000 \times 60000 \times 1/5000 = 60000\ tuples$$

$$NPages(Employee \bowtie Project) = 60000 / 100 = 600\ pages$$

Cost to join with Department = $NPages(Employee \bowtie Project) +$
 $NPages(Employee \bowtie Project) \times$
 $NPages(Department) - NPages(Employee \bowtie Project)$
 $= 200 \times 600 = 120000 \text{ I/O}$
 Overall cost = $30050 + 120000 = \mathbf{150050 \text{ I/O}}$

b.(plan2)

Cost to join Project \bowtie Department = $3 \times NPages(Project) + 3 \times$
 $NPages(Department)$
 $= 3 \times (600 + 200) = 2400 \text{ I/O}$
 Result size of Project \bowtie Department = $NTuples(Project) \times$
 $NTuples(Department) \times 1/NKeys(projid)$
 $= 60000 \times 20000 \times 1/60000 = 20000 \text{ tuples}$
 $NPages(Project \bowtie Department) = 20000 / 100 = 200 \text{ pages}$
 Cost to join with Employee = $2 \times NumPasses \times NPages(Project \bowtie$
 Department) + $2 \times NumPasses \times NPages(Employee) + NPages(Project \bowtie$
 Department) + $NPages(Employee) - NPages(Project \bowtie Department)$
 $= 4 \times 200 + 5 \times 50 = 1050 \text{ I/O}$
 Overall cost = $2400 + 1050 = \mathbf{3450 \text{ I/O}}$

b.(plan3)

Cost to select from Employee
 $= RF(\text{salary} < 300000) \times (NPages(\text{Index on salary}) + NPages(Employee))$
 $= 1/2 \times (10 + 50) = 30 \text{ I/O}$

Result size of $\sigma_{\text{salary} < 300000}(Employee) = NTuples(Employee) \times$

RF(salary < 300000)

= $5000 \times 1/2 = 2500$ tuples

NPages($\sigma_{\text{salary} < 300000}$ (Employee)) = $2500 / 100 = 25$ pages

Cost to join $\sigma_{\text{salary} < 300000}$ (Employee) \bowtie Project = NPages(salary < 300000 (Employee)) + NPages($\sigma_{\text{salary} < 300000}$ (Employee)) \times

NPages(Project) – NPages($\sigma_{\text{salary} < 300000}$ (Employee))

= $25 \times 600 = 15000$ I/O

Result size of $\sigma_{\text{salary} < 300000}$ (Employee) \bowtie Project

= NTuples($\sigma_{\text{salary} < 300000}$ (Employee)) \times NTuples(Project) \times
 $1/\text{NKeys}(\text{eid from joining table})$

= $2500 \times 60000 \times 1/2500 = 60000$ tuples

NPages($\sigma_{\text{salary} < 300000}$ (Employee) \bowtie Project)

= $60000 / 100 = 600$ pages

Cost to join with Department = $3 \times \text{NPages}(\sigma_{\text{salary} < 300000}$ (Employee) \bowtie Project) + $3 \times \text{NPages}(\text{Department}) - \text{NPages}(\sigma_{\text{salary} < 300000}$ (Employee) \bowtie Project)

= $2 \times 600 + 3 \times 200 = 1800$ I/O

Overall cost = $30 + 15000 + 1800 = \mathbf{16830}$ I/O

b.(plan4)

Cost to join Department \bowtie Project = $2 \times \text{NumPasses} \times$
 $\text{NPages}(\text{Department}) + \text{NPages}(\text{Department}) + (\text{NPages}(\text{Index on projid}) +$
 $\text{NPages}(\text{Project}))$

= $5 \times 200 + (200 + 600) = 1800$ I/O

Result size of Department \bowtie Project same as Plan 2 = 200 pages

Cost to join with Employee = $2 \times \text{NumPasses} \times \text{NPages}(\text{Department} \bowtie \text{Project}) + \text{NPages}(\text{Department} \bowtie \text{Project}) + \text{NPages}(\text{Employee})$

= $5 \times 200 + 50 = 1050$ I/O

Overall cost = $1800 + 1050 = \mathbf{2850}$ I/O